

Nguyễn Hoàng Anh

anhnguyendev.me

Email: anh487303@gmail.com

Mobile: +84-985-335-735

EDUCATION

- **Ho Chi Minh City Open University**

Bachelor of Science in Computer Science

Ho Chi Minh City, Vietnam

Aug 2023 – Present

PROJECTS

- **Personal Portfolio & Blog Platform** [Live](#) | [GitHub](#)

Aug 2025 – Feb 2026

- Architected a bilingual (EN/VI) portfolio and MDX blog using Next.js 16 App Router with Server Components, internationalization, and Incremental Static Regeneration (ISR).
- Built a production-ready contact pipeline using Next.js Server Actions with Zod validation, honeypot bot detection, configurable rate limiting, parallel email delivery, and non-blocking logging.
- Implemented a comprehensive SEO pipeline including dynamic metadata, sitemap, RSS feed, Open Graph images, JSON-LD structured data, and hreflang support for bilingual content.
- Optimized frontend performance using dynamic imports, LazyMotion, server-side syntax highlighting, and content-visibility to reduce client-side JavaScript and improve initial page rendering.

- **Amazing Project 2**

Feb 2016 – May 2017

- Began a group to create the first 36 hour student-run hackathon
- Led 50 staff to raise \$175,000 and managed relationships with sponsors

- **Amazing Project 3**

Jan 2015 – May 2015

- Assist in recruiting potential students to the CS program
- Attend info sessions and student Q&A events

SKILLS

- **Languages:** JavaScript, Golang
- **Frontend:** React, Nextjs
- **Backend:** Node.js (Express, NestJS), Gin
- **Databases:** PostgreSQL, MySQL, Redis
- **Tools:** Git, Docker, Linux, Kubernetes, Github Actions

CERTIFICATIONS

- **TOEIC Listening & Reading:** 945/990

ETS, 2026

EXPERIENCE

- **Google**

Software Engineer

Mountain View, CA

Oct 2016 – Present

- **TensorFlow:** TensorFlow is an open source software library for numerical computation using data flow graphs; primarily used for training deep learning models. Worked on APIs and performance for training models on Tensor Processing Units (TPU).
- **Apache Beam:** Apache Beam is a unified model for defining both batch and streaming data-parallel processing pipelines, as well as a set of language-specific SDKs for constructing pipelines and runners.

- **Coursera**

Senior Software Engineer

Mountain View, CA

Jan 2014 – Oct 2016

- **Notifications:** Service for sending email, push and in-app notifications. Involved in features such as delivery time optimization, tracking, queuing and A/B testing. Built an internal app to run batch campaigns for marketing, etc.
- **Nostos:** Bulk data processing and injection service from Hadoop to Cassandra and provides a thin REST layer on top for serving offline computed data online.
- **Workflows:** Dataduct an open source workflow framework to create and manage data pipelines leveraging reusable patterns to expedite developer productivity.
- **Data Collection:** Designed the internal survey and crowdsourcing platform which allowed for creating various tasks for crowdsourcing or embedding surveys across the Coursera platform.

- **Dev Environment:** Analytics environment based on Docker and AWS, standardized the Python and R dependencies. Wrote the core libraries that are shared by all data scientists.
- **Data Warehousing:** Setup, schema design and management of Amazon Redshift. Built an internal app for access to the data using a web interface. Dataduct integration for daily ETL injection into Redshift.
- **Recommendations:** Core service for all recommendation systems at Coursera, currently used on the homepage and throughout the content discovery process. Worked on both offline training and online serving.
- **Content Discovery:** Improved content discovery by building a new onboarding experience on Coursera. Used this to personalize the search and browse experience. Also worked on ranking and indexing improvements.
- **Course Dashboards:** Instructor dashboards and learner surveying tools, which helped instructors run their class better by providing data on Assignments and Learner Activity.

- **Lucena Research**

Data Scientist

Atlanta, GA

Summer 2012 and 2013

- **Portfolio Management:** Created models for portfolio hedging, portfolio optimization and price forecasting. Also created a strategy backtesting engine used for simulating and backtesting strategies.
- **QuantDesk:** Python backend for a web application used by hedge fund managers for portfolio management.

- **Georgia Institute of Technology**

Research and Teaching Assistant

Atlanta, GA

Jan 2012 – Dec 2013

- **Research Assistant – Machine Learning:** Research on machine learning for portfolio hedging and replication algorithms. Modeling low-risk & continuous-return strategies. Developed the Python library QSTK.
- **Teaching Assistant – Computational Investing:** The online course on Coursera, had more than 100,000 students enrolled. It was featured on the 11 Alive News and the Atlanta Journal Constitution. Involved in creating assignments, exams and conducting recitation sessions. Also taught the on-campus version of the course.